

3D Acoustic Wave Simulation with Multi-GPU NCCL Communication and MPI Group Parallelism

Copyright statements

This project uses **OpenMPI** for distributed-memory parallelism via MPI. OpenMPI is an open-source implementation of the Message Passing Interface, licensed under the New BSD License.

This project uses the **NVIDIA CUDA Toolkit** for GPU programming. CUDA is provided by NVIDIA under the NVIDIA Software License Agreement.

This project uses **NCCL** for fast multi-GPU communication. NCCL is distributed by NVIDIA under a permissive BSD-like license.

This project uses **libxml2** for XML configuration parsing. libxml2 is developed by the GNOME project and released under the MIT License.

This project uses **seggy.h** from the open-source Seismic Unix (SU) package developed by the Colorado School of Mines, which is released under a BSD-style license.

Project Structure

This project implements a 3D acoustic wave finite-difference simulation using CUDA and MPI. It leverages **NCCL** for communication between multiple GPUs in a domain decomposition scheme, and employs **MPI** for shot-domain parallelism.

- **main.cpp**: Main framework
 - **kernel.cu**: CUDA kernel implementations
 - **fd.h**: Custom header file
 - **seggy.h**: Segy header file
 - **Makefile**: Compilation instructions
 - **parameter.xml**: XML parameter card
 - **run.sh**: Sample launch script
-

Build Instructions

Before compiling, make sure you have the following:

- **CUDA Toolkit**
- **MPI**
- **NCCL**
- **libxml2**

Compile

```
make
```

This produces the executable `GPU3D`.

To clean up:

```
make clean
```

Parameters (XML)

All simulation parameters are read from an XML file named `parameter.xml`.

```
<parameter name="fm">10</parameter>      <!-- Source central frequency
-->
<parameter name="dy">10</parameter>      <!-- Grid spacing in y -->
<parameter name="dx">10</parameter>      <!-- Grid spacing in x -->
<parameter name="dz">10</parameter>      <!-- Grid spacing in z -->
<parameter name="pml">50</parameter>      <!-- PML thickness -->
<parameter name="ny">676</parameter>      <!-- Grid points in y -->
<parameter name="nx">676</parameter>      <!-- Grid points in x -->
<parameter name="nz">210</parameter>      <!-- Grid points in z -->
<parameter name="disx_shot_grid">10</parameter> <!-- Shot grid interval
in x -->
<parameter name="disy_shot_grid">10</parameter> <!-- Shot grid interval
in y -->
<parameter name="sourcex_min_grid">100</parameter> <!-- Minimum grid of
shot in x -->
<parameter name="sourcex_max_grid">100</parameter> <!-- Maximum grid of
shot in x -->
<parameter name="sourcey_min_grid">100</parameter> <!-- Minimum grid of
shot in y -->
<parameter name="sourcey_max_grid">100</parameter> <!-- Maximum grid of
shot in y -->
<parameter name="sz">0</parameter>        <!-- Shot depth -->
<parameter name="gz">0</parameter>        <!-- Receiver depth -->
<parameter name="scale_y">4</parameter>    <!-- Receiver grid interval
in y -->
<parameter name="scale_x">4</parameter>    <!-- Receiver grid interval
in x -->
<parameter name="gpu_num">8</parameter>    <!-- Number of GPUs per node
-->
<parameter name="id_of_group">4</parameter> <!-- Number of processes/GPUs
per process group -->
<parameter name="dt">5e-4</parameter>     <!-- Sampling interval -->
<parameter name="t">4.0</parameter>       <!-- Total sampling time -->
<parameter name="fvel"/>path/to/vel.bin</parameter> <!-- Velocity
model (binary file) -->
```

```
<parameter name="outfile">/path/to/out.segy</parameter>  <!-- Output SEG-  
Y filename -->
```

Ensure paths `fvel` and `outfile` are absolute or valid relative paths.

Run Instructions

To run the simulation:

```
bash run.sh
```

Or manually:

```
mpirun -np 9 --hostfile nodefile ./GPU3D parameter.xml
```

The number of processes is equal to the total number of GPU cards actually used plus one (as master process).

Output

The output is a **SEG-Y** file containing seismic records, written to the path defined by the `outfile` parameter.
