

3D Acoustic Wave Simulation with Multi-GPU NCCL Communication and MPI Group Parallelism

This project implements a 3D acoustic wave finite-difference simulation using CUDA and MPI. It leverages **NCCL** for communication between multiple GPUs in a domain decomposition scheme, and employs **MPI** for shot-domain parallelism.

Project Structure

- `main.cpp`: Main framework
 - `kernel.cu`: CUDA kernel implementations
 - `fd.h`: Custom header file
 - `segy.h`: Segy header file
 - `Makefile`: Compilation instructions
 - `parameter.xml`: XML parameter card
 - `run.sh`: Sample launch script
-

Build Instructions

Before compiling, make sure you have the following:

- **CUDA Toolkit**
- **MPI**
- **NCCL**
- **libxml2**

Compile

```
make
```

This produces the executable `GPU3D`.

To clean up:

```
make clean
```

Parameters (XML)

All simulation parameters are read from an XML file named `parameter.xml`.

```

<parameter name="fm">10</parameter>      <!-- Source central frequency
-->
<parameter name="dy">10</parameter>      <!-- Grid spacing in y -->
<parameter name="dx">10</parameter>      <!-- Grid spacing in x -->
<parameter name="dz">10</parameter>      <!-- Grid spacing in z -->
<parameter name="pml">50</parameter>      <!-- PML thickness -->
<parameter name="ny">676</parameter>      <!-- Grid points in y -->
<parameter name="nx">676</parameter>      <!-- Grid points in x -->
<parameter name="nz">210</parameter>      <!-- Grid points in z -->
<parameter name="disx_shot_grid">10</parameter> <!-- Shot grid interval
in x -->
<parameter name="disy_shot_grid">10</parameter> <!-- Shot grid interval
in y -->
<parameter name="sourcex_min_grid">100</parameter> <!-- Minimum grid of
shot in x -->
<parameter name="sourcex_max_grid">100</parameter> <!-- Maximum grid of
shot in x -->
<parameter name="sourcey_min_grid">100</parameter> <!-- Minimum grid of
shot in y -->
<parameter name="sourcey_max_grid">100</parameter> <!-- Maximum grid of
shot in y -->
<parameter name="sz">0</parameter>        <!-- Shot depth -->
<parameter name="gz">0</parameter>        <!-- Receiver depth -->
<parameter name="scale_y">4</parameter>    <!-- Receiver grid interval
in y -->
<parameter name="scale_x">4</parameter>    <!-- Receiver grid interval
in x -->
<parameter name="gpu_num">8</parameter>    <!-- Number of GPUs per node
-->
<parameter name="id_of_group">4</parameter> <!-- Number of processes/GPUs
per process group -->
<parameter name="dt">5e-4</parameter>     <!-- Sampling interval -->
<parameter name="t">4.0</parameter>       <!-- Total sampling time -->
<parameter name="fvel">/path/to/vel.bin</parameter> <!-- Velocity
model (binary file) -->
<parameter name="outfile">/path/to/out.segy</parameter> <!-- Output SEG-
Y filename -->

```

Ensure paths `fvel` and `outfile` are absolute or valid relative paths.

Run Instructions

To run the simulation:

```
bash run.sh
```

Or manually:

```
mpirun -np 9 --hostfile nodefile ./GPU3D parameter.xml
```

The number of processes is equal to the total number of GPU cards actually used plus one (as master process).

Output

The output is a **SEG-Y** file containing seismic records, written to the path defined by the **outfile** parameter.
